

Chen Fang

☎ (217) 979-0555

🌐 LinkedIn

✉ chenf3@illinois.edu

✉ Chen.Fang@nyulangone.org

📁 Portfolio

Education

University of Illinois Urbana-Champaign

Bachelor of Science in Mathematics and Computer Science

Dean's List

May 2027

GPA: 3.97/4.00

Major GPA: 4.00/4.00

Honors & Awards

Anthropic AI for Science — Rare Disease Research Grant

2026

Awarded \$50,000 in Claude credits; **1 of 100 projects selected from 750+ applications**

Proposal: *Zebra or Horse? A Phenopacket-Grounded Benchmark for Genodermatoses Versus Infectious Mimics*

Skills

Language & Tools: Unity, Unreal Engine, Blender, Python, PyTorch, C, C++, C#, Java, Docker, Git

Work Experience

NYU Langone Health

Adjunct Researcher

New York, NY

Aug. 2026 – Present

- Drafted and submitted the Anthropic AI for Science research proposal for Dr. Megan Coffee.
- Developing medical VLMs and reinforcement-learning post-training pipelines for rare dermatology diagnosis.

PediaMed AI

Member of Technical Staff

San Jose, CA

Jun. 2026 – Present

- Applying GazeAgent's query-driven analysis of children's gaze to in-home developmental screening, turning how children look and play into early developmental signals.
- Extending child-learning-inspired multimodal models (CVPR 2026 workshop) toward child-like continual learning in AI.

Research Experience

Rehg Lab, University of Illinois Urbana-Champaign

Research Intern, Advisor: James M. Rehg

Champaign, IL

Jan. 2026 – Present

- Leading MedXRAgent: **first-author** NAACL 2027 submission, ECCV 2026 workshop paper, CVPR 2026 **Spotlight** demo.

Health Care Engineering Systems Center, University of Illinois Urbana-Champaign

Research Intern, Advisor: Inki Kim

Champaign, IL

Jan. 2026 – May 2026

- Co-led GazeAgent, query-driven gaze analysis for AR smart glasses; IMWUT under review, CVPR 2026 **Spotlight** demo.

Secure Learning Lab, University of Illinois Urbana-Champaign

Machine Learning and AI Security Research Intern

Champaign, IL

Sep. 2025 – May 2026

- Co-led AudioGuard, an audio-LLM guardrail for speaker identification and content safety; ICLR 2027 under review.

Siebel School of Computing and Data Science, UIUC

Head Teaching Assistant (CS 521); Teaching Assistant and Mentor (CS 124)

Champaign, IL

Jan. 2024 – Dec. 2025

- Contributed to CS 521 exam and rubric design, grading, and course operations.
- Optimized the CS 124 online office hour queuing system and rules, reducing average student wait time by 9.3%.
- Led and trained a team of CS 124 mentees to become course assistants.

Projects

MedXRAgent (EgoMedAgent): Evidence-Grounded Egocentric Clinical Assistant

Jan. 2026 – Present

- Designed a perceive–reason–act agent that turns AR-glasses dialogue and egocentric video into a clinical vignette, retrieves 500 candidate diagnoses, and grounds the top five in web evidence (PubMed, MedlinePlus).
- Curated MedXR-Bench (795 training and 271 test egocentric video cases from NEJM Image Challenge) and post-trained Qwen3.5 with GRPO; MedXRAgent achieves the **best open-ended diagnosis score**.
- Built the CVPR 2026 **Spotlight** demo: a Unity AR-glasses client streaming speech and photos over WebSocket to MedASR and MedGemma 1.5 (4B), which draft clinical notes and advice with per-patient memory.

GazeAgent: Query-Driven Gaze Analysis on Smart Glasses

Jan. 2026 – Present

- Co-developed a pipeline that turns observer-worn smart-glasses video into a gaze timeline via promptable gaze estimation and scene grounding; a code-generating agent answers spoken queries about attention.
- **100% code-execution success and 86% clinical consistency** on 100 queries; significantly lower workload in a 16-participant user study (NASA-TLX, $p < 0.01$); IMWUT under review.

AudioGuard: Audio Safety Guardrail for Voice AI

Sep. 2025 – May 2026

- Co-developed SoundGuard (waveform-level detection of harmful sounds and speaker identities) and ContentGuard (ASR plus policy-grounded text moderation) for audio-input and audio-output safety.
- Built dataset loaders and evaluation for AudioSafetyBench (10K+ samples, 17 languages, 50+ speakers); higher accuracy at much lower latency than audio-LLM guardrails on five benchmarks.